

# HW3 experiment Report

---

Name: 李崇楷 Student ID: 313834006

Code: <https://github.com/kailee0422/RNN-Transformer/tree/main/HW3>

## Introduction

This project focuses on building a Named Entity Recognition (NER) tool using the BERT model. I train and evaluate the model on the DNRTI dataset, following standard procedures with train.txt, valid.txt, and test.txt. To further enhance performance, I also explore the use of SecBERT with BiLSTM and CRF. The goal is to understand and implement a modern NER pipeline based on transformer architectures.

## Method

This section details the implementation of our SecBERT–BiLSTM–CRF model for Named Entity Recognition (NER), covering data preparation, model design, embedding integration, training procedure, and evaluation.

### Data Preprocessing

- **Baseline tokenization** : For the baseline model, I use standard tokenization (e.g., spaCy or WordPiece from a generic BERT), replace digits with zero, lowercase all tokens, and pad/truncate sequences to a fixed length. No domain-specific normalization is applied.
- **SecBERT tokenization** : In addition to the baseline steps, I employ the SecBERT tokenizer with its security-optimized vocabulary. This tokenizer better handles security jargon and multi-word expressions common in vulnerability descriptions.
- **CoNLL-style Input**: I read raw files where each line contains a token and its IOB tag, sentences separated by blank lines. Malformed lines and lines with only one token are ignored.
- **Token–Tag Extraction**: For each sentence, we collect tokens and their corresponding tags into Python lists.
- **Tag Mappings**: We build tag2idx and idx2tag dictionaries by enumerating the sorted set of all tags across train/valid/test splits.
- **Sequence Encoding**: Using jackaduma/SecBERT tokenizer, we convert each sentence to subword token IDs (input\_ids) and attention masks, padded/truncated to a fixed length (256).
- **Label Alignment**: We align original word-level tags to subword positions, assigning the tag only to the first subword of each word and using the 'O' tag for padding or special tokens..

### Model Architecture:SecBERT–BiLSTM–CNN–CRF

My model begins by encoding each sentence with SecBERT (jackaduma/SecBERT), producing 768-dimensional contextualized embeddings with 0.1 dropout. A two-layer bidirectional LSTM (with 0.2 dropout) then transforms these embeddings into 512-dimensional token features (256 per direction). For character-level modeling when using CNN mode, we embed each character (25-dim), apply a 1×3 convolution and max-pool across characters

to yield a 25-dim vector, which is concatenated to the BiLSTM output. Finally, a linear layer projects the combined features to the tag space, and a CRF layer models tag transitions, optimizing negative log-likelihood during training and decoding via Viterbi at inference.

### Difference between Baseline and SecBERT

**Baseline Embedding:** the model uses either randomly initialized word embeddings or generic pre-trained embeddings (e.g. GloVe or standard BERT) without security tuning.

**SecBERT Embedding:** SecBERT is a BERT model further pre-trained on security-related corpora, producing contextual representations that capture domain-specific semantics (e.g. vulnerabilities, attack patterns).

The integration code remains identical; only the encoder weights differ, enabling direct ablation.

## Training Setup

### Baseline BiLSTM–CRF Training

- **Optimizer:** Stochastic Gradient Descent (SGD) with momentum = 0.9 and initial learning rate = 0.015.
- **Learning-rate Decay:** At the end of each epoch, decay LR by factor  $1/(1 + \text{decay\_rate} * \text{epoch})$  where  $\text{decay\_rate}=0.05$ .
- **Gradient Clipping:** Clip gradients to a maximum norm of 5.0 to prevent explosion.
- **Batching & Updates:** Online updates (batch size = 1 sentence) shuffled each epoch; evaluate on dev set after each full pass.
- **Epochs:** Up to 50 epochs, saving best checkpoint by dev F1; early stopping not applied explicitly.

### SecBERT–BiLSTM–CNN–CRF Training

- **Optimizer:** AdamW with weight decay ( $1e-4$ ) and initial learning rate =  $1e-4$ , better suited for fine-tuning large pre-trained Transformers.
- **Scheduler:** ReduceLROnPlateau monitoring validation F1; reduce LR by factor 0.5 if no improvement for 2 consecutive epochs.
- **Gradient Clipping:** Not applied (empirically unnecessary when fine-tuning SecBERT with AdamW).
- **Mini-batch Training:** Batch size = 8 sentences for efficient GPU utilization
- **Epochs & Checkpointing:** 10 epochs; save model checkpoint with highest validation F1.
- **Validation Frequency:** Evaluate on dev set at end of each epoch.

By differentiating optimizers, batch strategies, and learning-rate schedules, we tailor training to the characteristics of non-pretrained BiLSTM vs. large Transformer fine-tuning.


## Evaluation

We report both entity-level and overall performance:

Entity-level Precision: For each named-entity category (e.g., PER, LOC, ORG), we compute the precision of predicted spans.

Overall Metrics: We calculate micro-averaged precision, recall, and F1-score across all entity types to assess general model performance.

## Result



**Figure 1 .** Baseline(left) and After using SecBERT(right) B- and I-tag score comparison.

Entity	Precision	Recall	F1-score
B-Area	0.76	0.73	0.74
B-Exp	0.87	0.9	0.89
B-Features	0.78	0.51	0.62
B-HackOrg	0.68	0.53	0.6
B-Idus	0.45	0.56	0.5
B-OffAct	0.77	0.71	0.74
B-Org	0.65	0.61	0.63
B-Purp	0.68	0.31	0.43
B-SamFile	0.67	0.68	0.68
B-SecTeam	0.79	0.66	0.72
B-Time	0.89	0.89	0.89
B-Tool	0.59	0.46	0.52
B-Way	0.76	0.64	0.7
I-Area	0.5	0.5	0.5
I-Exp	1.0	0.89	0.94
I-Features	0.71	0.36	0.48
I-HackOrg	0.29	0.22	0.25
I-Idus	0.5	0.22	0.31
I-OffAct	0.65	0.52	0.58
I-Org	0.09	0.08	0.09
I-Purp	0.65	0.29	0.4
I-SamFile	0.71	0.91	0.8
I-SecTeam	0.25	0.22	0.24
I-Time	0.55	0.6	0.57
I-Tool	0.61	0.49	0.54
I-Way	0.73	0.58	0.65

Entity	Precision	Recall	F1-score
B-Area	0.86	0.92	0.89
B-Exp	0.98	0.99	0.99
B-Features	0.97	0.96	0.96
B-HackOrg	0.84	0.81	0.83
B-Idus	0.85	0.95	0.9
B-OffAct	0.84	0.83	0.84
B-Org	0.75	0.72	0.74
B-Purp	0.84	0.99	0.91
B-SamFile	0.98	0.77	0.86
B-SecTeam	0.98	0.92	0.95
B-Time	0.94	0.94	0.94
B-Tool	0.69	0.78	0.73
B-Way	0.96	0.97	0.97
I-Area	0.77	0.9	0.83
I-Exp	1.0	1.0	1.0
I-Features	0.99	0.91	0.94
I-HackOrg	0.77	0.77	0.77
I-Idus	0.69	1.0	0.82
I-OffAct	0.91	0.77	0.83
I-Org	0.75	0.76	0.76
I-Purp	0.81	1.0	0.9
I-SamFile	1.0	0.85	0.92
I-SecTeam	0.74	0.86	0.79
I-Time	0.98	0.84	0.9
I-Tool	0.69	0.78	0.73
I-Way	0.94	0.98	0.96

**Figure 2 .** Baseline(left) and After using SecBERT(right) B- and I-tag metrics comparison.

	Testing Precision	Testing Recall	Testing F1
Baseline	71.75%	56.52%	64.69%
SecBERT	<b>82.46%</b>	<b>84.50%</b>	<b>83.47%</b>

**Table 1.** Testing metrics with an Baseline and after using SecBERT for embedding.

## Conclusion

We compared a standard BiLSTM–CRF baseline with a SecBERT–BiLSTM–CNN–CRF model on the DNRTI NER task. Incorporating SecBERT raised test F1 from 64.7% to 83.5%, showing that domain-tuned embeddings greatly improve entity recognition in security text. This demonstrates the value of security-focused pre-training for specialized NER applications.